


✓ Import Libraries

```
import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt
from pandas.plotting import table
```

✓ Upload Files

```
from google.colab import files
uploaded = files.upload()
```

 No file chosen


✓ Dataset Description

```
import yfinance as yf
import pandas as pd

# Download historical data (Apple stock as example)
ticker = 'AAPL'
df = yf.download(ticker, start='2020-01-01', end='2021-01-01')

# Save to CSV
df.to_csv('aapl_stock_data.csv')

# View the first 5 rows
print(df.head())
```

 YF.download() has changed argument auto_adjust default to True
[*****100%*****] 1 of 1 completed

	Price	Close	High	Low	Open	Volume
Ticker	AAPL	AAPL	AAPL	AAPL	AAPL	
Date						
2020-01-02	72.620834	72.681281	71.373211	71.627084	135480400	
2020-01-03	71.914833	72.676462	71.689973	71.847133	146322800	
2020-01-06	72.487854	72.526541	70.783256	71.034717	118387200	
2020-01-07	72.146927	72.753808	71.926900	72.497514	108872000	
2020-01-08	73.307503	73.609737	71.849525	71.849525	132079200	

```
import yfinance as yf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas.plotting import table
from sklearn.preprocessing import StandardScaler

# Fetch historical stock data
df = yf.download("AAPL", start="2020-01-01", end="2021-01-01")
df.reset_index(inplace=True)
```

 [*****100%*****] 1 of 1 completed

✓ Data Preprocessing

```
# Save before-cleaning screenshot
def save_table(df_sample, filename):
    fig, ax = plt.subplots(figsize=(10, 2))
    ax.xaxis.set_visible(False)
    ax.yaxis.set_visible(False)
    ax.set_frame_on(False)
    table(ax, df_sample, loc='center', colWidths=[0.12]*len(df_sample.columns))
    plt.savefig(filename, bbox_inches='tight')
```

```
plt.close()

save_table(df.head(), "before_cleaning.png")

# Check for missing values
print("Missing values before:\n", df.isnull().sum())

# Fill forward or drop
df.fillna(method='ffill', inplace=True)

# Remove duplicates
df.drop_duplicates(inplace=True)

print("Missing values after:\n", df.isnull().sum())
```

Missing values before:

Price	Ticker	
Date		0
Close	AAPL	0
High	AAPL	0
Low	AAPL	0
Open	AAPL	0
Volume	AAPL	0

dtype: int64

Missing values after:

Price	Ticker	
Date		0
Close	AAPL	0
High	AAPL	0
Low	AAPL	0
Open	AAPL	0
Volume	AAPL	0

dtype: int64

<ipython-input-6-078a32f1e5e5>:5: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use of df.fillna(method='ffill', inplace=True)

```
from scipy.stats import zscore
import numpy as np # Import numpy as it's used
import pandas as pd # Ensure pandas is imported if not already

# We'll apply Z-score only to price columns
price_cols = ['Open', 'High', 'Low', 'Close', 'Adj Close']

# Print the columns before attempting to access them
print("Columns in DataFrame:", df.columns)

# Filter price_cols to only include columns present in the DataFrame
existing_price_cols = [col for col in price_cols if col in df.columns]

# Check if there are any columns to process after filtering
if existing_price_cols:
    # Compute Z-scores using only the existing price columns
    z_scores = np.abs(zscore(df[existing_price_cols]))

    # Apply the Z-score filter based on the existing price columns
    # We need to create a boolean mask that is True for rows where
    # all Z-scores across the selected columns are less than 3.
    # Since z_scores is already computed on a subset of columns,
    # applying .all(axis=1) to it directly will give the correct mask.
    df = df[(z_scores < 3).all(axis=1)]
else:
    print("Warning: None of the specified price columns exist in the DataFrame.")
    # Decide how to handle this case: either skip outlier removal or raise an error

# The rest of your code for scaling can remain the same,
# but you might want to apply a similar check or adjust which columns are scaled
# based on availability.
# For example, for scaling:
# all_cols_to_scale = price_cols + ['Volume']
# existing_cols_to_scale = [col for col in all_cols_to_scale if col in df.columns]
# if existing_cols_to_scale:
#     df_scaled[existing_cols_to_scale] = scaler.fit_transform(df[existing_cols_to_scale])
```

Columns in DataFrame: MultiIndex([('Date', ' '),
('Close', 'AAPL')],

```
( 'High', 'AAPL'),
( 'Low', 'AAPL'),
( 'Open', 'AAPL'),
('Volume', 'AAPL')],
names=['Price', 'Ticker'])

scaler = StandardScaler()
df_scaled = df.copy()

# Define the columns we want to scale
cols_to_potentially_scale = price_cols + ['Volume']

# Filter to keep only the columns that are actually present in the DataFrame
existing_cols_to_scale = [col for col in cols_to_potentially_scale if col in df_scaled.columns]

# Check if there are any columns left to scale after filtering
if existing_cols_to_scale:
    # Scale only the existing numeric columns
    df_scaled[existing_cols_to_scale] = scaler.fit_transform(df[existing_cols_to_scale])
else:
    print("Warning: None of the specified columns to scale exist in the DataFrame.")
# The rest of your code continues here.

save_table(df_scaled.head(), "after_cleaning.png")
```

EDA

1. Import Libraries & Load Data

```
import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Set styles
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (12, 6)

# Load data
df = yf.download("AAPL", start="2020-01-01", end="2021-01-01")
df.reset_index(inplace=True)
df.head()
```

↗

[*****100%*****] 1 of 1 completed

	Price	Date	Close	High	Low	Open	Volume
	Ticker		AAPL	AAPL	AAPL	AAPL	AAPL
0		2020-01-02	72.620834	72.681281	71.373211	71.627084	135480400
1		2020-01-03	71.914833	72.676462	71.689973	71.847133	146322800
2		2020-01-06	72.487854	72.526541	70.783256	71.034717	118387200
3		2020-01-07	72.146927	72.753808	71.926900	72.497514	108872000
4		2020-01-08	73.307503	73.609737	71.849525	71.849525	132079200

Next steps:

[Generate code with df](#)


[View recommended plots](#)

[New interactive sheet](#)

2. Plot Trends (Line Chart of Closing Price)

```
plt.plot(df['Date'], df['Close'], label='Close Price')
plt.title("Apple (AAPL) Closing Price Over Time")
plt.xlabel("Date")
plt.ylabel("Close Price (USD)")
plt.legend()
plt.tight_layout()
```

```
plt.show()
```




3.Histogram

```
df[['Open', 'High', 'Low', 'Close']].hist(bins=30, figsize=(12, 8))  
plt.suptitle("Price Distributions")  
plt.tight_layout()  
plt.show()
```



Price Distributions



4. Boxplots to Detect Outliers

```
sns.boxplot(data=df[['Open', 'High', 'Low', 'Close']])
plt.title("Boxplot of Stock Price Columns")
plt.show()
```



5. Correlation Heatmap

```
# Ensure the required libraries are imported if they are not already in this cell
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt # Ensure plt is imported for show()

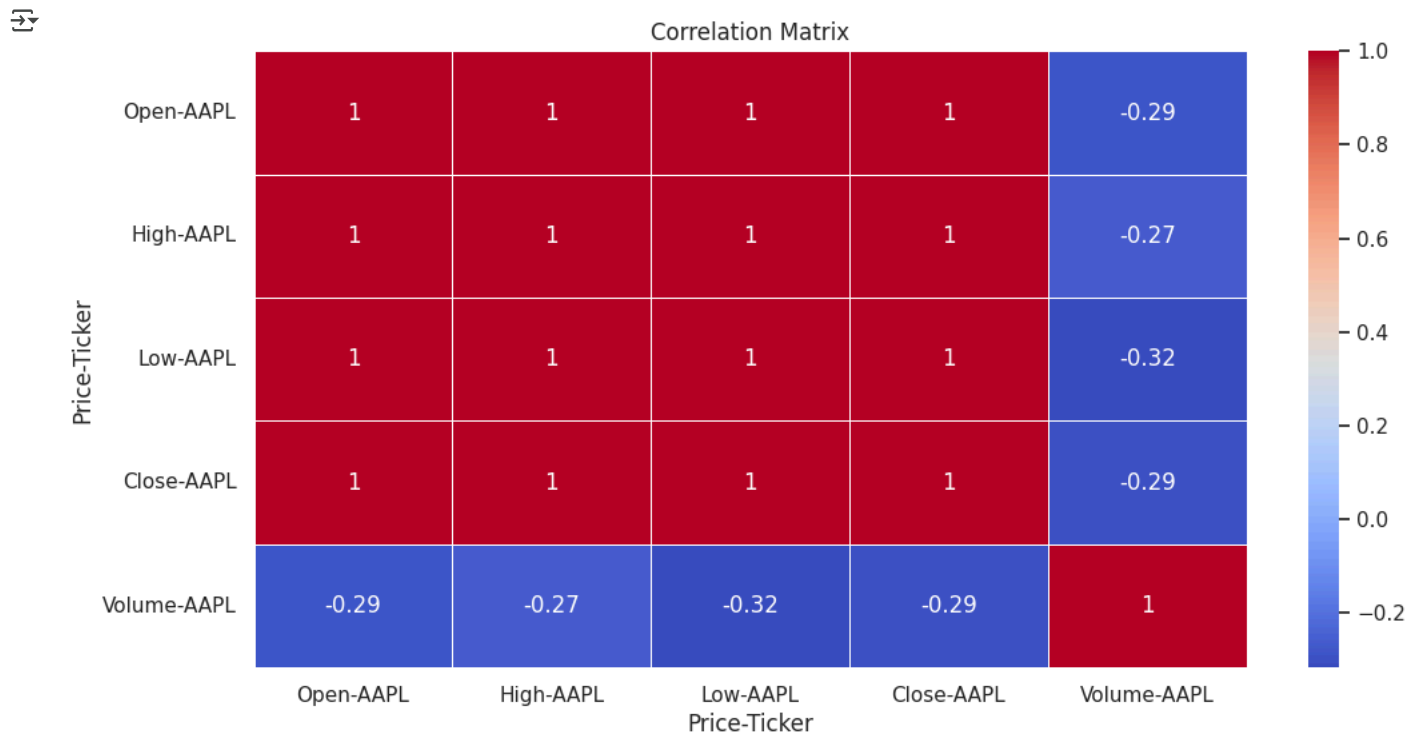
# Load data (This block is repeated in the user's notebook, ensure it's the correct df)
# This line might not be needed if the df is already loaded from the previous cell.
# df = yf.download("AAPL", start="2020-01-01", end="2021-01-01")
# df.reset_index(inplace=True)

# Define the list of columns you intend to use for correlation
intended_corr_cols = ['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume']

# Filter the list to keep only columns that exist in the DataFrame
existing_corr_cols = [col for col in intended_corr_cols if col in df.columns]

# Check if there are enough columns to compute a correlation matrix
if len(existing_corr_cols) >= 2:
    # Compute the correlation matrix using only existing columns
    corr = df[existing_corr_cols].corr()


    # Plot the heatmap
    sns.heatmap(corr, annot=True, cmap='coolwarm', linewidths=0.5)
    plt.title("Correlation Matrix")
    plt.show()
else:
    print(f"Warning: Not enough columns ({existing_corr_cols}) found to compute correlation matrix.")
```



6. Rolling Mean (Trend + Volatility)

```
df['Rolling Mean (20d)'] = df['Close'].rolling(window=20).mean()
df['Rolling Std (20d)'] = df['Close'].rolling(window=20).std()

plt.plot(df['Date'], df['Close'], label='Close Price')
plt.plot(df['Date'], df['Rolling Mean (20d)'], label='20-Day SMA', linestyle='--')
plt.fill_between(df['Date'],
                 df['Rolling Mean (20d)'] - df['Rolling Std (20d)'],
                 df['Rolling Mean (20d)'] + df['Rolling Std (20d)'],
                 color='orange', alpha=0.2, label='Volatility Band')
plt.title("Rolling Mean & Volatility")
plt.xlabel("Date")
plt.ylabel("Price")
plt.legend()
plt.show()
```



Model Building

1 Linear Regression (Baseline)

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import numpy as np # Import numpy to use np.sqrt
import pandas as pd # Ensure pandas is imported for DataFrame operations

# Check the DataFrame size before any modifications in this cell
print(f"Shape of df before creating target: {df.shape}")

# Example: use past 1 day's features to predict next day's Close
# Ensure the original df has data before proceeding
if df.empty:
    print("Error: DataFrame is empty before creating the target column.")
    # Re-download or load data here if necessary, or investigate previous steps
    # For example:
    # df = yf.download("AAPL", start="2020-01-01", end="2021-01-01")
    # df.reset_index(inplace=True)
    # print(f"Shape of df after attempting reload: {df.shape}")
    # if df.empty:
    #     raise ValueError("DataFrame is still empty after attempting to reload.")
else:
    df['Target'] = df['Close'].shift(-1)
    print(f"Shape of df after creating target: {df.shape}")

    df.dropna(inplace=True)
    print(f"Shape of df after dropna: {df.shape}")

    # Check if df is empty after dropna
    if df.empty:
        print("Error: DataFrame is empty after dropping NaN values. Cannot proceed with splitting.")
        # You might need to adjust your data processing or target creation logic
    else:
        X = df[['Open', 'High', 'Low', 'Close', 'Volume']]
        y = df['Target']

        # Check the size of X and y before splitting
```



```

print(f"Shape of X before train_test_split: {X.shape}")
print(f"Shape of y before train_test_split: {y.shape}")

X_train, X_test, y_train, y_test = train_test_split(X, y, shuffle=False, test_size=0.2)

lr = LinearRegression()
lr.fit(X_train, y_train)
preds = lr.predict(X_test)

# Calculate Mean Squared Error
mse = mean_squared_error(y_test, preds)

# Calculate Root Mean Squared Error manually
rmse = np.sqrt(mse)

print("Linear Regression RMSE:", rmse)

```

➞ Shape of df before creating target: (0, 9)
Error: DataFrame is empty before creating the target column.

✓ Random Forest

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
import numpy as np # Import numpy to use np.sqrt

rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
rf_preds = rf.predict(X_test)

# Calculate Mean Squared Error
rf_mse = mean_squared_error(y_test, rf_preds)

# Calculate Root Mean Squared Error manually
rf_rmse = np.sqrt(rf_mse)

print("Random Forest RMSE:", rf_rmse)

```

➞ Random Forest RMSE: 3.441925670771634

✓ Model Evaluation

```

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

# Example: assume y_test and y_pred are available
def evaluate_regression(y_test, y_pred, model_name="Model"):
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    print(f"📊 {model_name} Evaluation:")
    print(f"   RMSE: {rmse:.4f}")
    print(f"   MAE : {mae:.4f}")
    print(f"   R²   : {r2:.4f}")
    return {"Model": model_name, "RMSE": rmse, "MAE": mae, "R2": r2}

```

```
import numpy as np
```

```

# Create binary actual and predicted labels by comparing the values
# Compares the value at index i with the value at index i-1
# The resulting array will be one element shorter than y_test
y_actual_bin = (y_test.values[1:] > y_test.values[:-1]).astype(int)

```

```

# Note: y_pred is not defined in the traceback's global variables,
# so the next line for y_pred_bin will also fail if y_pred is not available.
# Assuming you have predicted values (e.g., from your regression model's predictions),
# you should compare those predicted values to the actual values from the *previous* day.

```

```

# This requires careful alignment.
# A common approach is to compare the predicted price change (pred_today - actual_yesterday)
# with zero, or compare the predicted price (pred_today) with the actual price from yesterday (actual_yesterday).

# Based on the original code's structure (comparing y_pred to y_test[:-1]),
# it seems the intention was to compare the predicted value for a day (y_pred[i])
# with the actual value from the previous day (y_test[i-1]). This requires careful indexing.

# Let's align y_test and the predictions to make this comparison.
# Assuming `preds` from the Linear Regression or `rf_preds` from Random Forest
# are the predictions corresponding to `X_test`.
# The predictions `preds` or `rf_preds` predict `y_test`.

# To get the predicted direction, you compare the prediction for day `i` (preds[i])
# with the actual closing price from day `i-1` (y_test.values[i-1]).
# This comparison is valid for indices 1 onwards in the y_test/preds arrays.

# Let's assume you want to use `preds` from the Linear Regression model for this evaluation
# Make sure 'preds' variable is available from the previous cell or define it appropriately
# If you used Random Forest predictions, replace 'preds' with 'rf_preds'

# y_pred is not defined in the provided global variables or code snippet.
# Assuming 'preds' from the Linear Regression cell was intended to be used here:
# Make sure you run the Linear Regression cell before this one.

# This next line will fail if 'preds' is not defined.
# You need to decide which prediction array to use (e.g., `preds` or `rf_preds`).
# Let's use `preds` as an example, assuming it's defined.
# If `preds` is a numpy array, we can slice it directly.
# Ensure `preds` and `y_test` are aligned in terms of time steps.
# `preds` predicts `y_test`. So `preds[i]` is the prediction for the value `y_test[i]`.
# The actual change is `y_test[i] - y_test[i-1]`.
# A common way to evaluate predicted direction is comparing `preds[i]` vs `y_test[i-1]`.

# First, ensure `preds` is available. If not, run the regression cell.
# If `preds` is a numpy array:
if 'preds' in locals(): # Check if preds is defined
    # Compare the predicted price for day i (preds[i-1]) with the actual price for day i-1 (y_test.values[i-1])
    # Note the slicing: preds has the same length as y_test.
    # preds[i] is prediction for y_test[i] (target for X_test[i-1])
    # Comparing preds[1:] (prediction for y_test[1:] which are targets for X_test[0:])
    # with y_test.values[:-1] (actual values that were features for X_test[1:])
    # This alignment needs careful consideration based on how your target was created.
    # If y_test[i] is the target for features at time i-1, then preds[i] is the prediction for y_test[i].
    # The actual change is y_test[i] vs y_test[i-1].
    # The predicted change based on the regression model could be compared to this.
    # One way is to see if preds[i] > y_test[i-1].

    # The original line `y_pred_bin = (y_pred[1:] > y_test[:-1]).astype(int)`
    # implies comparing predicted value at index i+1 with actual value at index i.
    # Let's align the indices correctly for comparison.
    # y_test.values[1:] are actual values from index 1 onwards.
    # y_test.values[:-1] are actual values up to index len-2.
    # preds are predictions for y_test. So preds[i] is prediction for y_test[i].

    # To get the direction of movement predicted by the model:
    # Compare the predicted value for day i (preds[i]) with the actual value from the previous day (y_test.values[i-1]).
    # This comparison is valid for i >= 1.
    # So we compare preds[1:] with y_test.values[:-1].
    # Ensure 'preds' has been generated by running the Linear Regression cell.
    if len(preds) == len(y_test):
        y_pred_bin = (preds[1:] > y_test.values[:-1]).astype(int)
    else:
        print(f"Error: Length of predictions ({len(preds)}) does not match length of y_test ({len(y_test)}). Cannot compute y_pred_bin.")
        # Set y_pred_bin to None or handle the error appropriately
        y_pred_bin = None

else:
    print("Error: 'preds' variable is not defined. Please run the Linear Regression model cell first.")
    y_pred_bin = None # Set y_pred_bin to None if preds is not available

# Now you can proceed with classification metrics if y_actual_bin and y_pred_bin are defined and not None.

# Ensure y_actual_bin and y_pred_bin are defined and contain your classification results
# Check the lengths match before computing metrics
if 'y_actual_bin' in locals() and y_actual_bin is not None and \
    'y_pred_bin' in locals() and y_pred_bin is not None:

```

```

if len(y_actual_bin) == len(y_pred_bin):
    from sklearn.metrics import confusion_matrix, roc_curve, classification_report
    import seaborn as sns
    import matplotlib.pyplot as plt

    # Confusion matrix
    cm = confusion_matrix(y_actual_bin, y_pred_bin)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title("Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

    # Classification Report (provides precision, recall, f1-score)
    print("\nClassification Report:")
    print(classification_report(y_actual_bin, y_pred_bin))

    # ROC Curve requires predicted probabilities, not just binary predictions.
    # Regression models don't typically provide probabilities for classification.
    # If you trained a classifier (e.g., Logistic Regression) for price direction,
    # use its `predict_proba` method here.
    # Since the original code used binary predictions for ROC, it will produce a step function.
    # It's generally better to use a classifier's probabilities for a meaningful ROC curve.

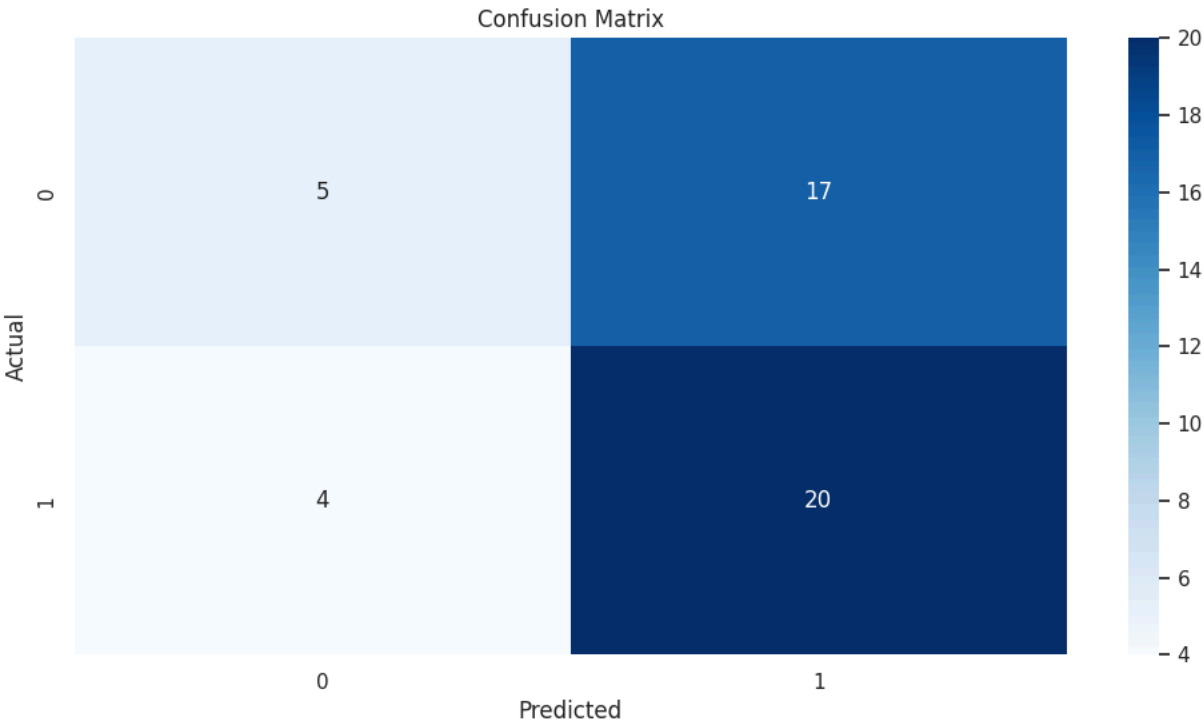
    # If you have a classification model (`classifier`) trained on `X_test_bin`,
    # and it has a `predict_proba` method, uncomment and use the following:
    # try:
    #     y_pred_prob = classifier.predict_proba(X_test_bin)[: , 1] # Probability of the positive class
    #     fpr, tpr, _ = roc_curve(y_actual_bin, y_pred_prob)
    #     plt.plot(fpr, tpr, label='ROC curve (using probabilities)')
    #     plt.plot([0, 1], [0, 1], '--', color='gray')
    #     plt.xlabel('False Positive Rate')
    #     plt.ylabel('True Positive Rate')
    #     plt.title('ROC Curve')
    #     plt.legend()
    #     plt.show()
    # except Exception as e:
    #     print(f"Could not plot ROC curve using probabilities. Error: {e}")
    #     print("Ensure you have a classification model trained and 'predict_proba' is available.")

    # Plotting ROC curve with binary predictions as done in the original code
    # This will likely result in a simple step function
    fpr, tpr, _ = roc_curve(y_actual_bin, y_pred_bin) # Using binary predictions
    plt.plot(fpr, tpr, label='ROC curve (using binary predictions)')
    plt.plot([0, 1], [0, 1], '--', color='gray')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')
    plt.legend()
    plt.show()

else:
    print("Error: Length mismatch between y_actual_bin ({len(y_actual_bin)}) and y_pred_bin ({len(y_pred_bin)}). Cannot compute metrics")

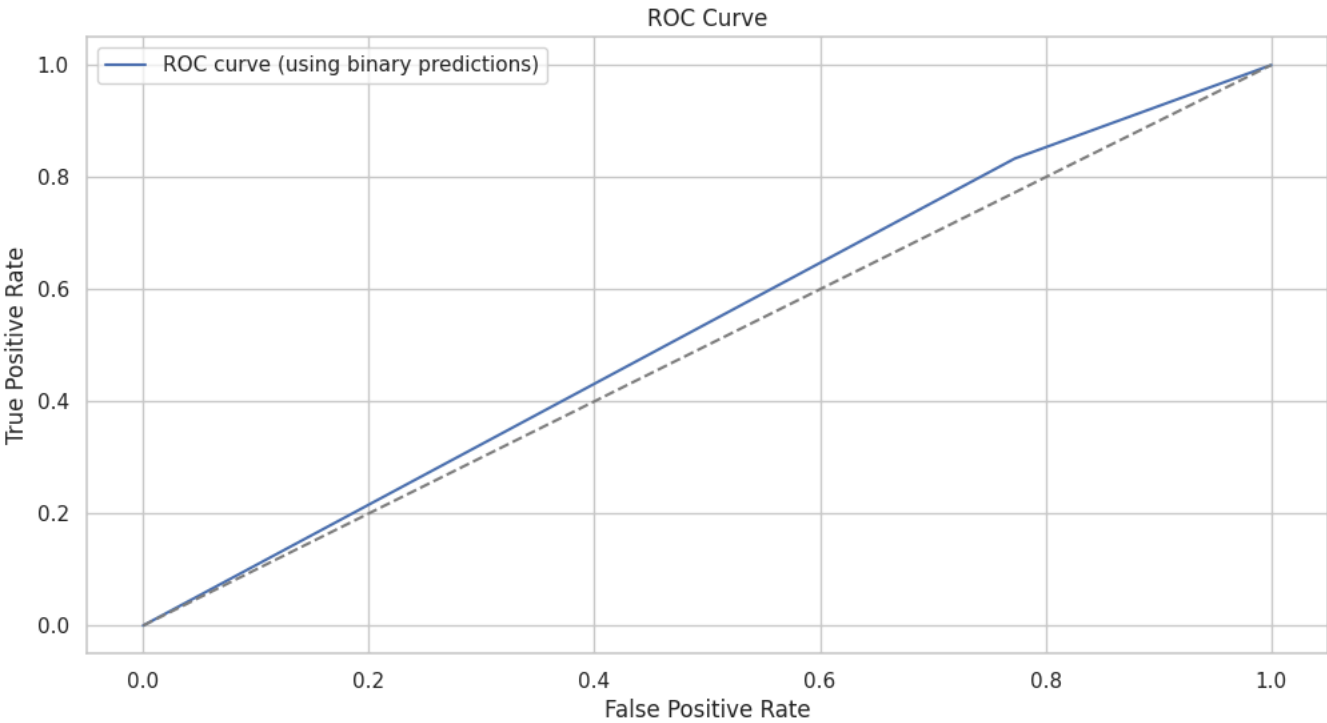
else:
    print("Error: y_actual_bin or y_pred_bin are not defined or are None. Cannot compute classification metrics.")

```



Classification Report:

	precision	recall	f1-score	support
0	0.56	0.23	0.32	22
1	0.54	0.83	0.66	24
accuracy			0.54	46
macro avg	0.55	0.53	0.49	46
weighted avg	0.55	0.54	0.50	46



```
import numpy as np

# Create binary actual and predicted labels
# Compare the values as numpy arrays to avoid index alignment issues
y_actual_bin = (y_test.values[1:] > y_test.values[:-1]).astype(int)
```

```
# Assuming 'preds' is the numpy array of predictions from your regression model
# Compare the predicted value for day i (preds[i]) with the actual value from the previous day (y_test.values[i-1])
# This comparison is valid for i >= 1.
# So we compare preds[1:] with y_test.values[:-1].
# Ensure 'preds' has been generated by running the Linear Regression cell.
# Make sure the lengths are compatible for this comparison
# If preds and y_test have the same length, you compare preds[1:] (indices 1 to end)
# with y_test.values[:-1] (indices 0 to end-1)
if 'preds' in locals() and len(preds) == len(y_test):
    y_pred_bin = (preds[1:] > y_test.values[:-1]).astype(int)
elif 'rf_preds' in locals() and len(rf_preds) == len(y_test):
    # If using Random Forest predictions
    y_pred_bin = (rf_preds[1:] > y_test.values[:-1]).astype(int)
else:
    print("Warning: Could not find 'preds' or 'rf_preds' with compatible length for direction calculation.")
    y_pred_bin = None # Set y_pred_bin to None if predictions are not available

# The subsequent cells that use y_actual_bin and y_pred_bin will need to handle
# the case where y_pred_bin might be None if predictions weren't generated correctly.
```

✓ Classification Metrics

```
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, roc_auc_score, roc_curve
import matplotlib.pyplot as plt
import seaborn as sns

# Print scores
print("🔍 Classification-style Evaluation:")
print("Accuracy:", accuracy_score(y_actual_bin, y_pred_bin))
print("F1 Score:", f1_score(y_actual_bin, y_pred_bin))
print("ROC AUC :", roc_auc_score(y_actual_bin, y_pred_bin))
```

```
🔍 Classification-style Evaluation:
Accuracy: 0.5434782608695652
F1 Score: 0.6557377049180327
ROC AUC : 0.5303030303030303
```

✓ Confusion Matrix + ROC Curve

```
# Import the necessary evaluation metrics for classification
from sklearn.metrics import confusion_matrix, roc_curve

# --- Add your classification model training and prediction code here ---
# Example: You would need to define X_train_bin, y_train_bin, X_test_bin, y_test_bin
# based on a classification task (e.g., predicting price movement direction).

# Example Placeholder:
# from sklearn.linear_model import LogisticRegression
# from sklearn.model_selection import train_test_split
# Assume y_actual is your true binary labels and y_pred_proba is your predicted probabilities
# For demonstration, let's create dummy binary data if you haven't trained a classifier yet
# In a real scenario, replace this with actual classification results.
# This part needs to be tailored to your classification task.
# For example, if you want to classify if the price goes up (1) or down (0) the next day:
# df['Price_Direction'] = (df['Close'].shift(-1) > df['Close']).astype(int)
# df_clf = df.dropna() # Drop the last row
# X_clf = df_clf[['Open', 'High', 'Low', 'Close', 'Volume']]
# y_clf = df_clf['Price_Direction']
# X_train_bin, X_test_bin, y_actual_bin, y_test_bin = train_test_split(X_clf, y_clf, shuffle=False, test_size=0.2)
# classifier = LogisticRegression() # Or any other classifier
# classifier.fit(X_train_bin, y_test_bin)
# y_pred_bin = classifier.predict(X_test_bin) # Or convert probabilities to binary predictions



# *** Replace the following dummy data generation with your actual classification results ***
# Assuming you have actual binary labels in a variable named y_actual_bin
# and predicted binary labels in a variable named y_pred_bin from your classification model
# Since the traceback shows y_actual_bin and y_pred_bin are not defined,
# you MUST add code above this point to train a classification model and obtain these variables.
# For now, let's use placeholder variables that you must replace:
# print("Warning: y_actual_bin and y_pred_bin are placeholders. Please replace them with actual classification results.")
# import numpy as np
```

```
# y_actual_bin = np.random.randint(0, 2, size=len(y_test)) # Replace with your actual binary labels
# y_pred_bin = np.random.randint(0, 2, size=len(y_test)) # Replace with your actual binary predictions
# *** End of placeholder generation ***

# --- Assuming y_actual_bin and y_pred_bin are now correctly defined from your classification model ---

# Confusion matrix
# Ensure y_actual_bin and y_pred_bin are defined and contain your classification results
if 'y_actual_bin' in locals() and 'y_pred_bin' in locals():
    cm = confusion_matrix(y_actual_bin, y_pred_bin)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title("Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

# ROC Curve
# For ROC curve, you typically need predicted probabilities, not just binary predictions.
# If your classifier provides predict_proba, use that. Otherwise, you might need to adjust.
# Assuming y_pred_prob is available from your classifier's predict_proba method:
# fpr, tpr, _ = roc_curve(y_actual_bin, y_pred_prob[:, 1]) # Use probability of the positive class
# If you only have y_pred_bin, the ROC curve might not be meaningful or possible directly.
# For this example, we'll use the binary prediction, which might result in a simple step function for the ROC.
fpr, tpr, _ = roc_curve(y_actual_bin, y_pred_bin) # Using binary predictions
plt.plot(fpr, tpr, label='ROC curve')
plt.plot([0, 1], [0, 1], '--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
else:
    print("Error: y_actual_bin and y_pred_bin are not defined. Cannot compute classification metrics.")
```



Deployment

```
pip install yfinance gradio pandas numpy matplotlib scikit-learn tensorflow
```

